


```
In [14]: import tensorflow as tf
from tensorflow.keras.datasets import mnist
from tensorflow.keras.utils import to_categorical
import matplotlib as plt
```

```
In [15]: (x_train, y_train), (x_test, y_test) = mnist.load_data()
```


```
In [20]: x_train = x_train.reshape(-1, 28, 28, 1)
x_test = x_test.reshape(-1, 28, 28, 1)
x_train = x_train.astype('float32') / 255
x_test = x_test.astype('float32') / 255

# Define a function to plot an image
def plot_image(image, figsize=(3, 3)): # Adjust figsize for desired image size
    plt.figure(figsize=figsize)
    plt.imshow(image.reshape(28, 28), cmap='gray')
    plt.axis('off')
    plt.show()


# Visualize some training images (modify range for more/Less images)
for i in range(9):
    plot_image(x_train[i])
    print(f"Label: {y_train[i]}")
```




Label: 5




Label: 0




Label: 4




Label: 1



Label: 9




Label: 2



Label: 1



Label: 3



Label: 1

```
In [3]: # Reshape for CNN input (28x28 pixels and 1 color channel)
```

```
x_train = x_train.reshape(-1, 28, 28, 1)
x_test = x_test.reshape(-1, 28, 28, 1)

# Normalize pixel values to [0, 1] range
x_train = x_train.astype('float32') / 255
x_test = x_test.astype('float32') / 255
```

```
In [4]: # Convert labels to one-hot vectors for multiclass classification
```

```
y_train = to_categorical(y_train)
y_test = to_categorical(y_test)
```

```
In [5]: model = tf.keras.Sequential([
    tf.keras.layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)),
    tf.keras.layers.MaxPooling2D((2, 2)),
    tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D((2, 2)),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(10, activation='softmax') # 10 output units for 10 digits
])
```

```
In [6]: model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
model.fit(x_train, y_train, epochs=10, batch_size=32, validation_data=(x_test, y_test))
```

```

Epoch 1/10
1875/1875 [=====] - 75s 37ms/step - loss: 0.1320 - accuracy: 0.9594 - val_loss: 0.0425 -
val_accuracy: 0.9867
Epoch 2/10
1875/1875 [=====] - 63s 34ms/step - loss: 0.0433 - accuracy: 0.9866 - val_loss: 0.0350 -
val_accuracy: 0.9884
Epoch 3/10
1875/1875 [=====] - 62s 33ms/step - loss: 0.0295 - accuracy: 0.9905 - val_loss: 0.0347 -
val_accuracy: 0.9881
Epoch 4/10
1875/1875 [=====] - 64s 34ms/step - loss: 0.0220 - accuracy: 0.9930 - val_loss: 0.0330 -
val_accuracy: 0.9898
Epoch 5/10
1875/1875 [=====] - 63s 33ms/step - loss: 0.0156 - accuracy: 0.9949 - val_loss: 0.0361 -
val_accuracy: 0.9882
Epoch 6/10
1875/1875 [=====] - 74s 40ms/step - loss: 0.0128 - accuracy: 0.9956 - val_loss: 0.0416 -
val_accuracy: 0.9882
Epoch 7/10
1875/1875 [=====] - 66s 35ms/step - loss: 0.0098 - accuracy: 0.9966 - val_loss: 0.0308 -
val_accuracy: 0.9911
Epoch 8/10
1875/1875 [=====] - 57s 31ms/step - loss: 0.0084 - accuracy: 0.9970 - val_loss: 0.0324 -
val_accuracy: 0.9908
Epoch 9/10
1875/1875 [=====] - 58s 31ms/step - loss: 0.0068 - accuracy: 0.9977 - val_loss: 0.0327 -
val_accuracy: 0.9920
Epoch 10/10
1875/1875 [=====] - 57s 31ms/step - loss: 0.0055 - accuracy: 0.9981 - val_loss: 0.0446 -
val_accuracy: 0.9904
<keras.callbacks.History at 0x254b748f0a0>

```

Out[6]:

In [7]: `y_pred = model.predict(x_test)`

```

313/313 [=====] - 3s 8ms/step

```

In [8]: `from sklearn.metrics import confusion_matrix`

```

cm = confusion_matrix(y_test.argmax(axis=1), y_pred.argmax(axis=1))
print("Confusion Matrix:")
print(cm)

```

Confusion Matrix:

```

[[ 978   0   0   0   1   0   0   1   0   0]
 [   0 1132   0   0   0   1   0   0   1   1]
 [   1   0 1019   0   2   0   0   8   2   0]
 [   0   1   0 980   0  24   0   2   3   0]
 [   0   0   0   0 976   0   1   0   2   3]
 [   0   0   0   1   0 889   1   0   1   0]
 [   2   2   0   0   2   4 946   0   2   0]
 [   0   5   0   0   0   0   0 1021   1   1]
 [   2   0   1   0   0   0   0   0 969   2]
 [   0   2   0   0   3   5   0   2   3 994]]

```

In [9]: `import numpy as np`

```

accuracy = np.mean(y_test.argmax(axis=1) == y_pred.argmax(axis=1))
print("Accuracy:", accuracy)

```

Accuracy: 0.9904

In [10]: `# Make predictions on a small subset of test data`

```

predictions = model.predict(x_test[:10])

```

`# Get the predicted class for each sample (argmax gives the index of maximum value)`

```

predicted_classes = predictions.argmax(axis=1)

```

`# Get the actual labels from the one-hot encoded test data`

```

actual_labels = y_test[:10].argmax(axis=1)

```

`# Print the predicted and actual labels for the first 10 samples`

```

print("Sample Predictions and Actual Labels:")

```

```

for i in range(10):

```

```

    print(f"Sample {i+1}: Predicted: {predicted_classes[i]}, Actual: {actual_labels[i]}")

```

1/1 [=====] - 0s 86ms/step
Sample Predictions and Actual Labels:
Sample 1: Predicted: 7, Actual: 7
Sample 2: Predicted: 2, Actual: 2
Sample 3: Predicted: 1, Actual: 1
Sample 4: Predicted: 0, Actual: 0
Sample 5: Predicted: 4, Actual: 4
Sample 6: Predicted: 1, Actual: 1
Sample 7: Predicted: 4, Actual: 4
Sample 8: Predicted: 9, Actual: 9
Sample 9: Predicted: 5, Actual: 5
Sample 10: Predicted: 9, Actual: 9